

# Lab 2 Report

## Description of Functions

k\_tsk\_create

k\_tsk\_exit

k\_tsk\_set\_prio

k\_tsk\_get

Scheduler

find\_value

compare

Swap

left/right child

Parents

Increase key

Decrease key

Insert

ExtractMax

Maximum

Reset priority

RemoveID

## Mandatory Testing Scenarios

Creating and terminating user-mode tasks

Sanity test

Create and destroy multiple tasks

Scheduling tasks with the same priorities

Scheduling tasks with different priorities

Sanity tests

Ability to deallocate memory that is owned

No ability to deallocate memory that is not owned

Getting correct information on tasks

## Extra Testing Scenarios

tsk create failure tests

tsk get failure tests

Set priority tests

Set prio failure tests

tsk failures at max task bounds

TCB reusal test

# Description of Functions

## **k\_tsk\_create**

1. We start k\_tsk\_create by error checking against cases such as
  - If the given task pointer is null
  - If the task entry pointer is null
  - If we have already reached the max number of tasks
  - If the stack size to create is too small
  - If the stack size to create is too big
  - If the stack size is not 8 byte aligned
  - If the given priority is invalid
2. After the initial error checks, we search for a free TCB
3. We then set up the the appropriate task information such as priority, stack size, entry or if its privileged
4. Call k\_tsk\_create\_new
5. Increment the number of active tasks

## **k\_tsk\_exit**

1. Set the current task's state to be 'dormant'
2. Set the current task to what the scheduler feels is appropriate
3. Deallocate the previously running tasks stack
4. Decrement the number of currently running tasks
5. Pop that task id off the heap we use as a priority queue
6. Run an error check to see if we should actually keep running the old task
7. Set the state of the current task to 'running' if needed
8. Switch to the old stack if needed

## **k\_tsk\_set\_prio**

1. Run checks for the following cases
  1. Invalid priority
  2. if the given task is already dormant
  3. if the current tasks priority is the same we can return early
  4. If the task id is valid
2. If the current task is privileged, let it set whatever priority it wants
3. If the given task is privileged, error
4. otherwise set the priority as it was requested
5. Update our heap(which acts as a priority queue) with the newly set priority
6. Call k\_tsk\_yield so that a new task can run if needed

## **k\_tsk\_get**

1. Run checks for the following cases
  1. If the given buffer is null
  2. if the given taskID maps to a dormant task
  3. If the given taskID is valid
2. populate the buffer with the correct information from the global list of TCBs
  1. this includes setting kernal and user stack pointers

## **Scheduler**

Our underlying data structure that handles the scheduling of tasks is a priority queue. Under the hood, this is implemented using a binary heap. This lets us pop off the highest priority runnable task in  $O(1)$  while also being able to insert items in  $O(\log n)$ .

The scheduler function extracts the TID off the top of the heap (priority queue). It also checks if the current task should keep running. In the case that a new task

needs to be run, it removes the next runnable task from the heap and then inserts the currently running task.

There are a few helper methods that we use to make working with the heap easier to manage:

### **find\_value**

- This method iterates through the heap to find a given task id
- this operation runs in  $O(n)$  but given that the heap is at most, `MAX_TASKS(16)`, the linear time search is acceptable

### **compare**

- a simplified comparison to compare the priority of two TCBs that get stored in the heap
- this is used when "heapifying"

### **Swap**

- Swaps two tcb's pointers

### **left/right child**

- returns the left or right child of a parent using pointer arithmetic to save on using space on pointers

### **Parents**

- returns the parents of a node using pointer arithmetic to save on using space on pointers

### **Increase key**

- iteratively moves a TCB up the heap until it is in the correct position

### **Decrease key**

- Iteratively pushes a TCB deeper into the heap until it is in the correct position

## **Insert**

- inserts a TCB into the heap in the correct position

## **ExtractMax**

- fetches the highest priority runnable task

## **Maximum**

- Extracts the TID of the highest priority runnable task

## **Reset priority**

- sets the priority of a given TCB and then moves it to the correct position in the heap

## **RemoveID**

- removes a given TCB from the heap

---

# **Mandatory Testing Scenarios**

## **Creating and terminating user-mode tasks**

### **Sanity test**

1. Create task A with a variety of stack sizes
2. let the task run and use the stack
3. Call tsk exit

### **Create and destroy multiple tasks**

1. Create tasks until the max number of tasks are created

2. get tasks for all TIDs to validate that get works as expected
3. assert that calling tsk\_create fails
4. tsk\_exit
5. tsk\_create should re-use the TID

## **Scheduling tasks with the same priorities**

1. Create tsks all with the same priority
2. Pop off the tasks using the scheduler
3. Assert that the popped off tasks are in the same order they were created

## **Scheduling tasks with different priorities**

### **Sanity tests**

1. assert using PRIO\_NULL fails
2. assert using an invalid TID fails
3. assert set prio on dormant task fails
4. assert set prio on kernel tasks fails
5. assert a simple test with valid set prio passes and can then have tsk\_get be called

## **Ability to deallocate memory that is owned**

1. Create a task that owns data and calls malloc
2. From that task, try to dealloc the data
3. assert success

## **No ability to deallocate memory that is not owned**

1. Create a task that owns data and calls malloc
2. create a task that try to dealloc the first task's data

3. assert that dealloc fails

## Getting correct information on tasks

1. Create a task
  2. call get on the task and verify in memory that the correct task was fetched
  3. validate that stack size, task id, priority, state, privilege, stack pointers are correct
  4. repeat the above for multiple tasks created at once
- 

## Extra Testing Scenarios

### tsk create failure tests

1. Create a task with a stack size too small
2. assert failure
3. create task with memory that is not 8 byte aligned
4. assert failure
5. create tasks with invalid priority input
6. assert failure
7. create tasks with null priority
8. assert failure
9. create tasks with invalid tid
10. assert failure
11. create tasks with invalid task function
12. assert failure

## **tsk get failure tests**

1. create task with ID 0
2. assert failure
3. use a null buffer
4. assert failure
5. use a task id that is out of bounds
6. assert failure
7. use a task that is dormant
8. assert failure

## **Set priority tests**

1. set a task prioity as low
2. get the task
3. assert the priority was set correctly
4. set task priority to high
5. get the task
6. assert the priority was set correctly

## **Set prio failure tests**

1. create task with ID 0
2. assert failure
3. use a null buffer
4. assert failure
5. use a task id that is out of bounds



6. assert failure
7. use a task that is dormant
8. assert failure
9. use an invalid priority
10. assert failure

## **tsk failures at max task bounds**

1. create tasks until at max\_tasks bound
2. assert next creation fails

## **TCB reusal test**

1. create task
2. yield task
3. creates task using same tid
4. assert that task id reused
5. tsk exit